

# Running and Managing Pods

---



**Anthony E. Nocentino**

ENTERPRISE ARCHITECT @ CENTINO SYSTEMS

@nocentino [www.centinosystems.com](http://www.centinosystems.com)

# Course Overview



**Using the Kubernetes API**

**Managing Objects with Labels, Annotations,  
and Namespaces**

**Running and Managing Pods**

# Overview

**Understanding Pods**

**Controllers and Pods**

**Multi-container Pods**

**Managing Pod Health with Probes**

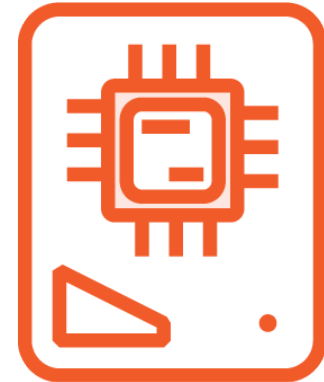
# What Is a Pod?



**Wrapper around  
your container  
based application**



**One or more  
containers**



**Resources**

# What Is a Pod? (con't)



**Unit of scheduling**

**Allocating work**

**A process that's running in your cluster**

**Unit of deployment**

**Your application configuration**

**Resources - Networking and Storage**

# Why do we need Pods?

Provide higher level abstraction  
over a container for manageability

# How Pods Manage Containers



**Single Container Pods**



**Multi-Container Pods**

# Single Container Pods



**Most common deployment scenario**

**Generally a single process running in a container**

**Often leads to easier application scaling**



# Controllers and Pods



**Controllers keep your apps in the desired state**

**Responsible for starting and stopping Pods**


**Application scaling**

**Application recovery**

**You don't want to run bare/naked Pods**

**They won't be recreated in the event of a failure**

# Working with Pods - kubectl



```
kubectl exec -it POD1 --container CONTAINER1 -- /bin/bash  
kubectl logs POD1 --container CONTAINER1  
kubectl port-forward pod POD1 LOCALPORT:CONTAINERPORT
```

# Demo

## Running Pods

- Bare Pods
- Creating Pods in a Deployment
- Using port-forward to access a Pod's application

# Multi-container Pods



**Tightly coupled applications**

**Scheduling processes together**

**Requirement on some shared resource**

**Usually something generating data while the other process consumes**

**Don't use this to influence scheduling, we use other techniques for that!**

# Multi-container Pods

```
apiVersion: v1
kind: Pod
metadata:
  name: multicontainer-pod
spec:
  containers:
  - name: nginx
    image: nginx
    ports:
    - containerPort: 80
  ...
  - name: alpine
    image: alpine
```


# Common Anti-Pattern for Multi-container Pods




Recovery Options

Limits Scaling

# Shares Resources Inside a Pod




**Networking**



**Storage**

# Shared Resources Inside a Pod - Networking




**Shared loopback interface, used for communication over localhost**

**Be mindful of application port conflicts**



# Shared Resources Inside a Pod - Storage



**Each container image has it's own file system**

**Volumes are defined at the Pod level**

**Shared amongst the containers in a Pod**

**Mounted into the containers' file system**

**Common way for containers to exchange data**


# Demo

**Running Multi-container Pods**

**Sharing data between containers in a Pod**

# Pod Lifecycle

Creation



Administratively

Controller

Running



Scheduled to a Node

Termination



Process is terminated/crashed

Pod is deleted

Evicted due to lack of resources

Node failure or maintenance

No Pod is redeployed

# Stopping/Terminating Pods

Grace Period Timer  
(30 sec default)

Pods changes to  
Terminating

**SIGTERM**

Service Endpoints  
and Controllers  
updated

IF > Grace Period  
**SIGKILL**

API and etcd are  
updated

kubectl delete pod  --grace-period=<seconds>

**Force Deletion - Immediately deletes records in API and etcd**

kubectl delete pod --grace-period=0 --force

# Persistence of Pods




**A Pod is never redeployed**

**If a Pod stops, a new one is created based on its Controller**

**Go back to the original container image(s) in the Pod definition**

# Persistence of Pods



**Configuration is managed externally**

**Pod Manifests, secrets and ConfigMaps**

**Passing environment variables into containers**

**Data Persistence is managed externally**

`PersistentVolume`

`PersistentVolumeClaim`

# Container Restart Policy



**A container in a Pod can restart independent of the Pod**

**Applies to containers inside a Pod and defined inside the Pod's Spec**

**The Pod is the environment the container runs in**

**Not rescheduled to another Node, but restarted by the Kubelet on that Node**

**Restarts with an exponential backoff, 10s, 20s, 40s capped at 5m and reset to 0 after 10m of successful runtime**

# Container Restart Policy



**Always (default) - will restart all containers inside a Pod**

**OnFailure - Non-graceful termination**

**Never**



```
apiVersion: v1
kind: Pod
metadata:
  name: nginx-pod
spec:
  containers:
  - name: nginx
    image: nginx
  restartPolicy: OnFailure
```

## Pod with Container Restart Policy

```
kubectl apply -f nginx.yaml
```

# Demo

**Pod lifecycle**

**Killing a container process**

**Container Restart Policy**

# Defining Pod Health



**A Pod is considered ready when all containers are ready**

**But we'd like to be able to understand a little more about our applications**

**We can add additional intelligence to our Pod's state and health**

**Container Probes**

`livenessProbe`

`readinessProbe`

# livenessProbes



**Runs a diagnostic check on a container**

**Per container setting**

**On failure, the Kubelet restarts the container**

**Container Restart Policy**

**Give Kubernetes a better understanding of our application**

# readinessProbes



**Runs a diagnostic check on the container**

**Per container setting**

**On startup, your application won't receive traffic until ready**

**On failure, removes Pod from load balancing or replication controller**

**Applications have long startup times**

**Prevents users from seeing errors**

# Types of Diagnostic Checks for Probes

Exec	tcpSocket	httpGet
Process exit code	Successfully Open a Port	Return Code 200 => and < 400
Success	Failure	Unknown

# Configuring Container Probes



**initialDelaySeconds** - number of seconds after the container has started before running container probes

**periodSeconds** - probe interval, default 10 seconds

**timeoutSeconds** Probe timeout 1 seconds

**failureThreshold** - number of missed checks before reporting failure, default 3

**successThreshold** - number of probes to be considered successful and live, default 1

```
spec:
  containers:
    ...
    livenessProbe:
      tcpSocket:
        port: 8080
      initialDelaySeconds: 15
      periodSeconds: 20
    readinessProbe:
      tcpSocket:
        port: 8080
      initialDelaySeconds: 5
      periodSeconds: 10
```

livenessProbes and readinessProbes



# Demo

## Implementing container probes

- `livenessProbes`
- `readinessProbes`

# Summary

**Understanding Pods**

**Controllers and Pods**

**Multi-container Pods**

**Managing Pod Health with Probes**

Thank You!

@nocentino

[www.centinosystems.com](http://www.centinosystems.com)